

PlayerOne

An AI playtester. It plays your game, catches the bug, and files the report with proof.

A free local AI does the playing. **A big AI only does the judging.**

`roaring-alfajores-8a941a.netlify.app`
`github.com/nijathatamli/AI-hakaton`



Small studios ship bugs because nobody has time to replay the game

Indie teams of 1 to 15 people build PC games in Godot, Unity and Unreal with no QA team. The first bug report is often a one-star review with no steps to reproduce.

The outcome

A bug report after every build: error, stack trace, repro steps, the frames before it, and a clip.

Manual testing

Slow and repetitive. Nobody replays level 3 for the hundredth time after every build.

Hiring QA

Too expensive for a small team, and still slow to turn around.

A big AI playing

Reads a full screenshot every move: 67,000 to 85,000 tokens per minute of play.

Split the brain. Cheap hands, smart judge.

1

Play

Free, on your PC

A local vision AI or our free bot plays the game with real keys and mouse. Any engine, no SDK.

>

2

Watch

Exact, zero tokens

Reads the engine console, catches crashes, hangs and soft-locks, keeps the last 12 seconds.

>

3

Judge

Your AI, tiny digests

Claude, Codex or Gemini read a 45 KB summary, decide what is a real bug and write the ticket.

The prototype


Pick a game. It knows how to play.

PlayerOne recognises the game, suggests the controls, and you approve exactly what the AI may press. It never touches Alt+F4 or anything outside the game. F8 stops it anywhere.

v0.2.3 released with installers

5 targets: Windows x64, x86, ARM64, macOS, Linux

MCP Claude Code, Codex, Gemini CLI direct it



The desktop app: pick a window, approve controls, press Start.

What the AI actually does, and what it does not

Plays from pixels

A vision model reads each frame and picks keys, mouse look and clicks from the approved controls. Stuck? It turns away and tries something else.

Works out the controls

Recognises the game by name or from the screen and picks a layout: first-person, platformer, top-down.

Judges with evidence

The big AI gets console errors and a contact sheet of the seconds before each incident, decides if it is a bug, and steers the player.

Detection is rules, on purpose

Crash, hang, console error and soft-lock checks need no model. That is why the clean build had 0 false alarms.

What a report looks like

Proof, not "it broke somewhere"

Error + stack trace

```
res://main.gd:154
```

Repro steps


the recorded inputs that led there

Contact sheet

36 KB, the seconds before

Clip

480p H.264, 14 KB



Contact sheet: 9 frames before the crash, 36 KB.

We broke our own game on purpose

Five planted bugs plus a clean build. One minute of real play each, three players, 18 runs. Every log, frame and clip is in the repo.

Bug	Random keys	Free bot	Local AI 3B
Soft-lock	34 s	15 s	39 s
Script error	missed	12.5 s	missed
Crash	missed	18 s	missed
Freeze	missed	20 s	missed
Coin counter	missed	missed	missed
False alarms	0	0	0

What broke, and what we did

- PlayerOne froze with a frozen game. Capture now has a 1.5 s deadline.
- 13 false alarms. It now tries every control before calling it stuck. Down to 0.
- 8 unit tests now cover the console classifier and token meter.

Compared with today

The same minute of testing, 96% fewer tokens

70,026

tokens a big AI spends playing one minute itself

about \$0.14 per minute

2,975

tokens it spent judging that minute through PlayerOne

about \$0.006 per minute

12–20 s

to catch a crash, a hang or a script error, report already written

vs a human replaying the level

Token counts from the meter in each session.json. Dollar figures at Claude Sonnet 5.5 input price, \$2 per million tokens.

Failures we have not fixed

What PlayerOne still gets wrong

The coin counter bug

No error, no stall. Only a model that understands the screen can see it. The 3B model did not.

The local AI is slow

25 to 45 moves a minute against 77 for the bot. The bot is the better driver; the AI is better at judging.

A weak judge writes weak reports

The 3B judge wrote "press a key". PlayerOne now falls back to the recorded inputs.

Small sample

One run per bug, on bugs we planted ourselves. Next: real third-party games.

Runs today, on hardware people already own

Data needed

None. No training data and no game source code. It works from the screen and the engine console.

Running cost

Free on the user's PC. With a cloud judge, about \$0.36 per hour of play instead of \$8.40.

Already live

Accounts, plans and usage metering on Supabase. Stripe webhook. Installers for 5 targets.

Who pays

Free

Local AI, crash and console capture. Costs us nothing to serve.

Connect

Monthly. Their own Claude, Codex or Gemini key judges.

Studio

Per seat. Shared reports, a playtest in CI after every build.

Credits

Our hosted judge, priced on the tokens we meter.

Originality and next step

No SDK. No engine plugin. Any game on screen.

Other AI testing tools need engine integration and run the expensive model on every frame. PlayerOne attaches to any window, trusts the engine console first, and only sends the big model a compressed digest.

Try it

`roaring-alfajores-8a941a.netlify.app`
Download, pick your game, press Start.

Next 30 days

Pilot with three Baku indie studios on their real games. Measure bugs found per build.

Then: Unity and Unreal log probes on real projects, and signed builds.